

Assignment 6 DESIGN.pdf

Ananya Batra

March 12, 2023

1 Description

Purpose: Write a program to perform LZ78 compression on any text or binary file and a program to decompress any file compressed with the previous algorithm, respectively. Both programs operate on big and little endian systems.

Structure: A user can compress a file with **encode.c**, which implements LZ78 compression. LZ78 compression is a lossless compression that compresses a file by storing a sequence of bytes (words) into a dictionary that gets built during the compression process. Words are denoted by codes in the dictionary. The compression algorithm works by replacing any future occurrence of a word that exists in a dictionary with its code. These words are stored in a Trie ADT, implemented by **trie.c**. Each TrieNode contains a symbol and n child nodes (256 in our case). In addition to **trie.c**, **encode.c** also calls upon functions in **io.c**. For this assignment, both encode and decode programs perform read and writes in blocks of 4KB. **io.c** contains functions in order to facilitate efficient reading and writing. A user can decompress a file (compressed with **encode.c**) with **decode.c**. Decompression works by translating codes to words (using a lookup table called a Word Table which is just an array of words). The Word ADT is implemented in **word.c**. Makefile compiles all the c program files and builds the encode and decode executables.

2 Design

trie.c

trie_node_create(): TrieNode constructor; creates a node with the code set by the parameter code

- dynamically allocate memory for a TrieNode node
- set node's code = code
- loop through all the node's children and set them to null
- return node

trie_node_delete(): TrieNode destructor; frees all the memory allocated for a TrieNode

- if the node is not null, free the memory allocated for the node and set it equal to NULL

trie_create(): Initializes a trie by creating a root TrieNode with code EMPTY_CODE

- create root with trie_node_create() with code EMPTY_CODE
- if the root pointer is NULL, return NULL
- otherwise, return root

trie_reset(): resets a trie to only contain the root node

- if this doesn't work, loop through the root's children and call trie_delete on them
- set each child = NULL

trie_delete(): deletes a subtree (starting at the node n passed in as a parameter)

- if the node isn't null

- loop through the node's children
- call `trie_delete()` on each of the children
- call `trie_node_delete` on `n`

trie_step(): returns a pointer to the child node of a node with the value `sym`

- return `n`'s child at index `n`

word.c

word_create(): Constructor for Word

- dynamically initialize memory for a Word `word`
- sets `word`'s `len = len`
- dynamically allocate memory for `word`'s `syms`. Create dynamically initialized array with `len` number of elements with each element being the size of a `uint8_t`.
- return `word`

word_append_sym(): Constructs a new Word which is the same as the Word passed in appended with `sym`

- dynamically initialize memory for a Word `new_word`
- set `new_word`'s `len = w->len + 1`
- loop through `w->len`
 - `new_word`'s `sym` at every index = `w ->syms[i]`
- set `new_word -> syms` at index `w->len = sym`
- return `new_word`

word_delete(): frees all the memory allocated to a word

- free the memory allocated to `w`'s `syms`

- free w

wt_create(): Creates a new word table

- initialize table as an array of word pointers
- if table is null, return null
- set the word at index EMPTY_CODE to an empty string of length 0 with word_create()
- return table

wt_reset(): resets a word table to only contain the empty word

- loop through all indices until MAX_CODE with i
 - if i is not the EMPTY_CODE, delete the word at index i with word_delete() and set the word at index i to null

io.c

read_bytes(): Helper function to perform reads. This function repeatedly calls read() until the number of bytes specified are read or there are no more bytes left to read.

- initialize a variable total_bytes_read to keep track of all the bytes read and a variable bytes_read_in() to keep track of the bytes read in at each call of read()
- while total_bytes_read is less than the bytes to be read in
 - use read() to read in the bytes specified by to_read - total_bytes_read already from the infile into the buffer.
 - add the number of bytes read in to bytes_read_in
 - if bytes_read_in is 0, return the bytes_read_in
 - add bytes_read_in to total_bytes_read
- return bytes_read_in

write_bytes(): Helper function to perform writes. This function repeatedly calls `write()` until the number of bytes specified are written out or there are no more bytes left to write.

- initialize a variable `total_bytes_written` to keep track of all the bytes written out and a variable `bytes_written()` to keep track of the bytes written out at each call of `write()`
- while `total_bytes_written` is less than the bytes to be written (`to_write`)
 - use `write()` to write out the bytes specified by `to_write - total_bytes_written` to the outfile from the buffer
 - add the number of bytes written to `bytes_written`
 - add `bytes_written` to `total_bytes_written`
- return `total_bytes_written`

read_header(): Populates the specified header from infile

- use `read_bytes()` to populate header from infile
- if the bytes are in `big_endian` ordering
 - use `swap32()` to swap the bytes for header's magic to make it little endian
 - use `swap16()` to swap the bytes for header's protection to make it little endian
- if header's magic is not the magic number specified (AKA the file was not compressed by the compression algorithm described later in this document), exit the program

write_header(): Writes out the header to outfile

- if the file is in `big_endian` ordering
 - use `swap32()` to swap the bytes for header's magic to make it little endian

- use `swap16()` to swap the bytes for header's protection to make it little endian
- write the bytes to outfile

read_sym(): Sets `sym` = to the current sym in the buffer and returns whether there are any more symbols left to read

- initialize a variable `end_of_buf = 0`
- if this is the first time reading in (aka `max_sym_index` which is a global static variable is negative one), then set `max_sym_index =` to bytes read in with `read_bytes()` and set `curr_sym_index = 0`
- set `sym =` to the value at `sym_curr_index` in `sym_buffer`
- to figure out what to return:
- if `sym_curr_index > max_sym_index - 1`, increment `sym_curr_index` and return true because there's still more indices in the buffer to read in
- otherwise, if `max_sym_index = BLOCK`, there are a couple cases to take care of
 - set the buffer to 0 with `mem_set` and set `sym_curr_index = 0`
 - read in `BLOCK` bytes into `sym_buffer` from `infile` and set the bytes read equal to `max_sym_index`.
 - if `max_sym_index` is 0, return false because this means we're at the end of the buffer.
 - otherwise, return true because there are still more bytes to be read in
- otherwise if the current index is at the last index in the buffer and `end_of_buf = 0` (this is for the case where the buffer is not full AKA we're at our last read)
 - increment `sym_curr_index`
 - increment `end_of_buf`
 - return true

- otherwise return false

write_pair(): a pair consists of a code followed by a symbol. This function sets the current bytes in the bit buffer to equal the code and symbol so that they can be written out when the buffer gets full

- loop through the bits of code from $i = 0$ to $i = \text{bitlen} - 1$
 - if the current index is greater than $\text{BLOCK} * 8$ (block is in bytes - multiplying it with 8 gets the number of bits in block)
 - * write the buffer to outfile with `flush_pairs()`
 - check if the current bit in code is 1. The way to do this is the right shift code by i bits (to get the i th bit in the lowest bit position) and ANDing it with 1. If the result is 1, then the i th bit in code is a 1.
 - * OR the `bit_buffer[bit_curr_index/8]` (to get the byte) with (1 left shifted by $(\text{bit_curr_index} \% 8)$). This sets the byte's current index with respect to the current byte to 1
 - increment `bit_curr_index`
- do the same process for adding sym to the bit_buffer. This time, loop from $i = 0$ to 7 because sym is only 1 byte

flush_pairs(): writes out the bit buffer and resets it

- first initialize a variable called `curr_index`
- if `bit_curr_index` is divisible by 8, set `curr_index = bit_curr_index/8`. otherwise, set `curr_index = bit_curr_index/8 + 1`. This is because C does floor division. So when figuring out which byte `curr_index` is at, we need to add 1 to the result of `bit_curr_index/8`.
- write out the bytes from bit_buffer to outfile
- set add the bytes written out $* 8$ to `total_bits`
- set `bit_curr_index` to 0

- reset the buffer with `memset()`

read_pair(): sets the code and symbol by reading in the values into the bit buffer

- initialize code and sym as 0 since we only plan to be setting the bits that are 1
- if `max_bit_index` (global static variable initialized to -1) is -1 (AKA this is the first time we're reading from the `bit_buffer`)
 - read BLOCK bytes from `bit_buffer` and set `max_bit_index` to $8 * (\text{bytes read in})$ to get the number of bits read in
 - add the number of bits to `total_bits`
- loop from `i = 0` to `bitlen - 1`
 - if `bit_curr_index` is equal to `max_bit_index` (which means we've gone through all the bytes in the `bit_buffer` already)
 - * reset the buffer and set the `bit_curr_index = 0`
 - * read in more bytes and set the number of bytes = $\text{max_bit_index} / 8$
 - * add the number of bits to `total_bits`
 - check if the current bit in `bit_buffer`'s current byte is 1. The way to do this is the right shift `bit_buffer[bit_curr_index / 8]` by $(\text{bit_curr_index} \% 8)$ bits (to get the current bit in the lowest bit position) and ANDing it with 1. If the result is 1, then the current bit in the `bit_buffer` is a 1.
 - * if it is, OR `*code` with 1 left shifted by `i`
 - increment `bit_curr_index` by 1
- do the same process to set `*sym`
- return true if code is not `STOP_CODE` and false otherwise

write_word(): sets the current byte in the symbol buffer to be the current word and writes out the buffer when it gets full

- loop through the length of word with i
- if the current index of the buffer goes beyond the end of the buffer size (4096), write out the bytes in the buffer.
- the idea behind this is we want to fill the buffer with the symbols in word. Once the buffer is full, write out all the bytes
- set the element at the current index of the buffer = w-i syms[i]
- increment current index of the buffer

flush_words(): writes out the sym buffer to outfile

- write the bytes from sym_buffer to outfile and add it to total_syms
- set sym_curr_index = 0
- reset sym_buffer

encode.c - main: contains the encode program as described at the top of this document

- parse through the commandline options using switch statements
 - 'i' opens the file specified to read the message to compress from
 - 'o' opens the file specified to write the compressed message to
 - 'v' sets the statistics flag to true
 - 'h' sets the help flag to true and will later help print help message
- create a stat struct called buffer
- obtain information about infile and store it into buffer with fstat
- dynamically initialize memory for a FileHeader pointer
- set header's magic and protection (get protection using st_mode)
- set the outfile's permission to match header's protections

- write the header to outfile
- Now, to compress (based on the pseudocode provided in the assignment sheet):
- create a root node and create a copy of it called `curr_node`. create a `prev_node` and `next_code` and initialize both to NULL
- initialize integers `curr_sym`, `prev_sym`, and `next_code` (initialized to `START_CODE`)
- Loop through all the symbols in infile with `read_sym()`
 - if `next_node` is not null (meaning current prefix has been encountered before), set `prev_node = curr_node` and `curr_node = next_node`
 - otherwise
 - * write the pair to outfile
 - * let `curr_node`'s child at index `curr_sym` be a new `trie_node` whose code is `next_code`
 - * set `curr_node = root`
 - * increment `next_code` by 1
 - if `next_code` is the `MAX_CODE`
 - * reset the trie
 - * set `curr_node = root`
 - * set `next_code = START_CODE`
 - set `prev_sym = curr_sym`
- free the header and close infile and outfile

decode.c - main: contains the decode program as described at the top of this document

- parse through the commandline options using switch statements
 - 'i' opens the file specified to read the compressed message from
 - 'o' opens the file specified to write the decompressed message to

- 'v' sets the statistics flag to true
 - 'h' sets the help flag to true and will later help print help message
- initialize a FileHeader header and read the header from infile into the header
- set the permissions for the outfile as the header's protections
- create a new word table
- initialize curr_sym, curr_code, and next code (initialized to START_CODE)
- loop through all the pairs with read_pair()
 - table[next_code] = append the read symbol to the word denoted by the read code
 - write the word at index next_code in the word table to the outfile
 - if next_code = MAX_CODE, reset the word table and set next_code = START_CODE
- flush any buffered words to outfile
- if statistics is enabled
 - set uncompressed_size = total_syms and initialize compressed_size to 0
 - if total_bits is divisible by 8, set compressed_size = total_bits / 8 + sizeof(FileHeader). otherwise, set compressed_size = total_bits / 8 + sizeof(FileHeader) + 1. This is because C does floor division. So when figuring out which byte total_bits is at, we need to add 1 to the result of total_bits/8 if it's not a perfect multiple of 8 (to round up).
 - calculate the percent difference with the formula:
 $(1 - \text{compressed_size} / \text{uncompressed_size}) * 100$
 - print all three values to stderr
- delete the word table
- close infile and outfile

3 Credit

- I went to Yiyuan's section on 2/27 to understand the compression algorithm
- I went to the last 30 minutes of Audrey's group tutoring session on 2/27 to understand io.c
- I went to Eric's tutoring session on 2/28 to clear up some misconceptions I had about trie.c and on 3/2 to get some insight on read_pair()
- I went to Jessie's tutoring session on 3/6. She went over io functions and the discussion helped me fix my read_sym()
- I went to Miles' tutoring session on 3/7. He helped me fix my write_pair() and write_word(). Essentially, I was using the wrong index to set the byte in the buffer.
- I went to John's session on 3/7. He went over read_pair() and write_pair(). He helped me debug my read_pair() too.
- I utilized the pseudocode provided in the assignment sheet to write encode.c and decode.c